

A Whitepaper

Reducing RTT of DNS Query Resolution using RFC 7706

*Dr. Balaji Rajendran, Gopinath Palaniappan, Sanjay Adiwal,
Shubham Goyal, Bindhumadhava B S*

January 2020



Centre of Excellence in

DNS
SECURITY

Background

Domain Name System (DNS) is the backbone of the Internet. The DNS service has considerable stake in ensuring genuine and quick Internet availability. The Recursive Resolvers (RR) are an important component of the DNS infrastructure. Most of the DNS query resolutions happen at the RR (its cache) itself without the need to visit all the components in the DNS resolution hierarchy, while some happen after visiting all components of the DNS resolution hierarchy. Some RR servers face the issue of lengthier round-trip-time (RTT) to their closest DNS root server, which happens to be the first step for any DNS query resolution. In this document, we discuss on cutting down the RTT between the RR and its closest DNS root server by implementing RFC 7706 and analyze the overall change in the time taken for DNS query resolution.

Approach

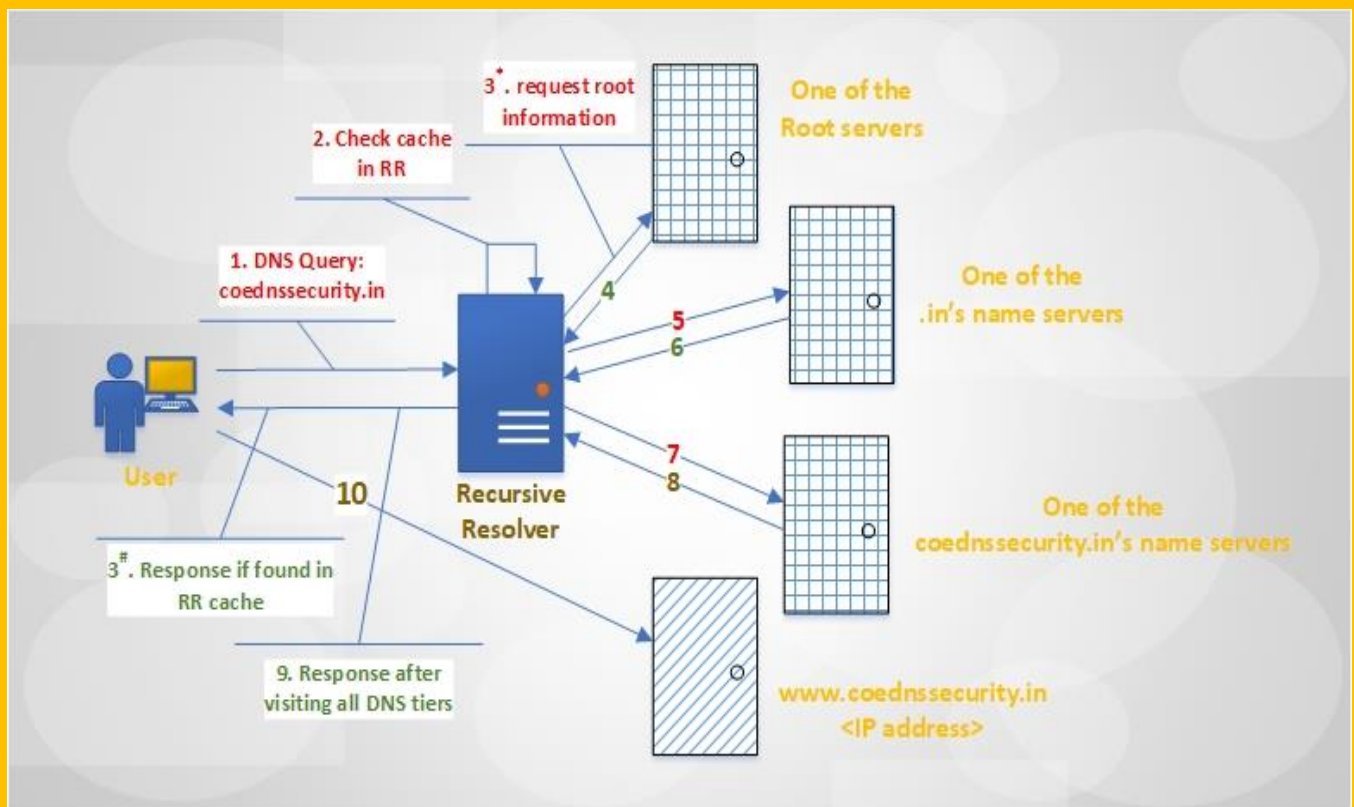


Figure.1: Conventional approach of DNS query resolution

In the **conventional approach**, the RR at the first place tries to resolve a DNS query by checking for the domain name in its own cache, if not found in its cache, it queries the closest root server, and follows it with subsequent queries to the TLD and SLD servers. This conventional approach is represented in Figure.1 above. The domain resolution happens along the steps 1, 2 & 3[#] [[#]resolved from RR cache] or 1, 2, 3^{*}, 4, 5, 6, 7, 8 & 9 [^{*}RR cache doesn't have the domain name and so proceeds for fetching data from root servers] as represented in Figure.1 above, depending on whether the RR server is able to resolve the DNS query from its cache or not.

“In the conventional approach, the RR server spends considerable time to reach out to the closest root server”

In our **proposed approach**, as per **RFC 7706** [1], we recommend creating an up-to-date root zone server on a loopback address on the same host as that of the RR server, and use that loopback server when the RR server has to look up for root information. The RR server would validate all responses received from the root server on the loopback address, just as it would do for all responses from a remote root server.

Our proposed approach is represented in Figure.2 below. In this proposed approach, the domain name resolution happens along the steps 1, 2 & 3[#] [[#]resolved from RR cache] or 1, 2, 3^{*}, 4, 5, 6, 7 & 8 [^{*}RR cache doesn't have the domain name and so fetches root data from its own loopback server], depending on whether the RR server is able to resolve the DNS query from its cache or from its loopback server containing root data. One can notice that the hosting of loopback server has reduced the number of steps from 9 to 8, meaning visiting the root servers to fetch root data relevant to the DNS query has been dropped.

“In the proposed approach, the loopback server containing root data is hosted on the RR server itself to avoid visiting the root servers for root data”

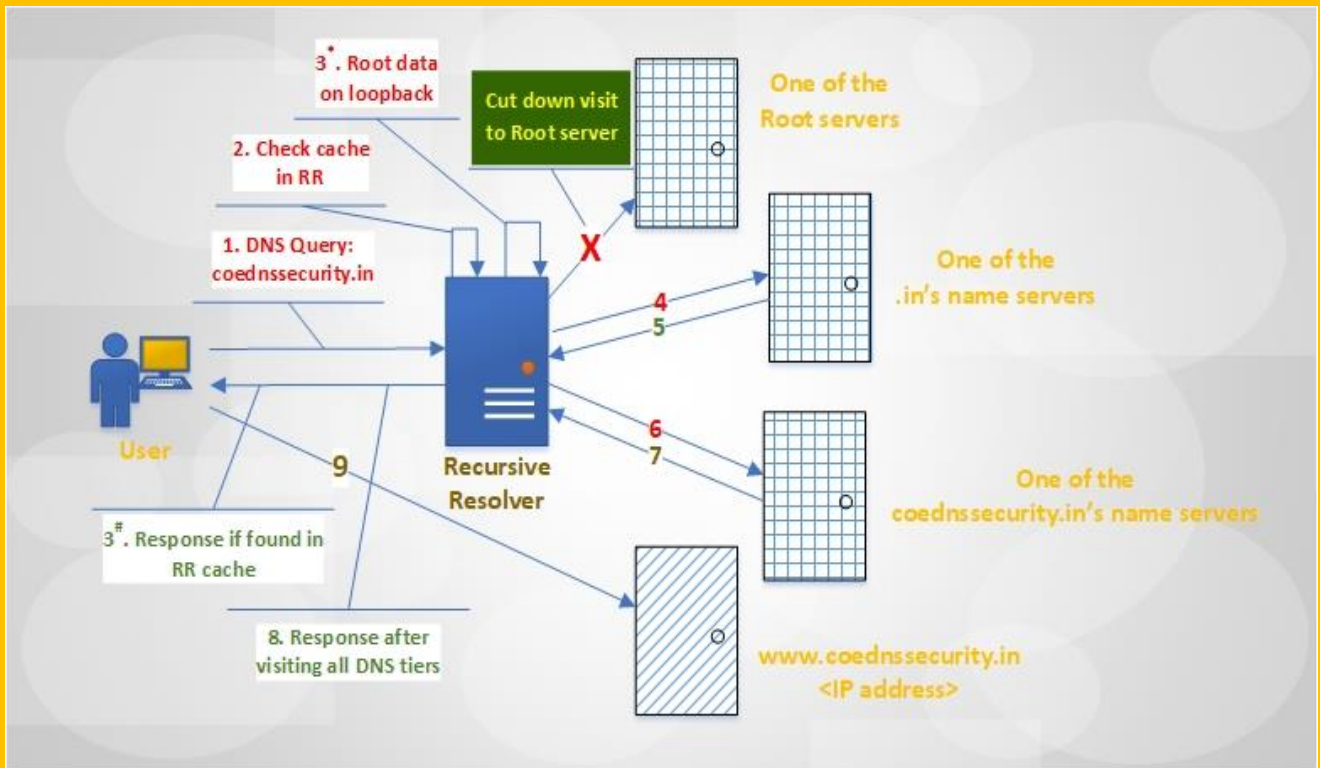


Figure.2: Proposed approach comprising of RFC 7706 implementation

Setting-up root zone at RR loopback

Pre-requisites:

- (a) The RR server must be able to run an authoritative server on one of its loopback addresses (127/8 in IPv4 and ::1 in IPv6).
- (b) The RR server must be able to retrieve a copy of the entire root zone including all DNSSEC-related records, and validate them.
- (c) The RR Server must have a software to do the configuration changes. We use Bind 9.9.4 [2] for implementation here.

Procedure:

- (a) Get the root zone data with all required DNSSEC records needed for validation. The root zone data should be obtained by AXFR over TCP from anyone of the following mentioned in the Table.1 below:

Root server operators	b.root-servers.net
	c.root-servers.net
	f.root-servers.net
	g.root-servers.net
	k.root-servers.net
DNS servers from ICANN	xfr.lax.dns.icann.org
	xfr.cjr.dns.icann.org

Table.1: Root Zone Data

(b) To configure the Recursive Resolver (RR) on Bind 9.9.4 [2] for implementing a root zone on its loopback address, we need to list the following in its configuration file:

```
view root {
    match-destinations { 127.0.0.1; };
    zone "." {
        type slave;
        file "rootzone.db";
        notify no;
        masters {
            192.228.79.201;           # b.root-servers.net
            192.33.4.12;             # c.root-servers.net
            192.5.5.241;             # f.root-servers.net
            192.112.36.4;            # g.root-servers.net
            193.0.14.129;            # k.root-servers.net
            192.0.47.132;            # xfr.cjr.dns.icann.org
            192.0.32.132;            # xfr.lax.dns.icann.org
            2001:500:84::b;           # b.root-servers.net
            2001:500:2f::f;          # f.root-servers.net
            2001:7fd::1;             # k.root-servers.net
            2620:0:2830:202::132     # xfr.cjr.dns.icann.org
            2620:0:2d0:202::132     # xfr.lax.dns.icann.org
        };
    };
};
```

```
view recursive {  
    dnssec-validation auto;  
    allow-recursion { any; };  
    recursion yes;  
    zone "." {  
        type static-stub;  
        server-addresses { 127.0.0.1; };  
    };  
};
```

Our Set-up

We have set-up an **Open, Public Recursive Resolver** for both IPv4 and IPv6 traffic and made it available for Internet users Worldwide, which is configured on Bind 9.9.4 [2] with RFC 7706 implemented in it and hosted at IPv4 address: **223.31.121.171** and IPv6 address: **2405:8a00:8001::20**.

Our experimental set-up comprised of a dedicated terminal with a public IP address in order to query the RR. We followed the procedure as mentioned below, first with RFC 7706 and next without RFC 7706 on the RR:

- a) We clear the RR cache before querying.
- b) We downloaded 100 distinct domain names from IANA, and created 5 files of 20 domains in each.
- c) We implemented a linux shell script which accepts an input of 20 domain names and queries the RR with them and writes the RTT on to a file upon receiving the response.



Figure.3: Procedure followed in the experiment

Analysis of RTT with and without RFC 7706 implementation

The results of our experiment are listed in the table below and also represented as a graph in Figure.3. Our experiment shows that the RTT for domain name resolution was clearly lesser in the RR with the RFC 7706 implementation.

Description	Time in milliseconds	
	with RFC 7706	without RFC 7706
Minimum RTT	106	138
Maximum RTT	4094	4869
Average RTT	1224.89	1621.89

Table.2: RTT for domain resolution with and without RFC 7706

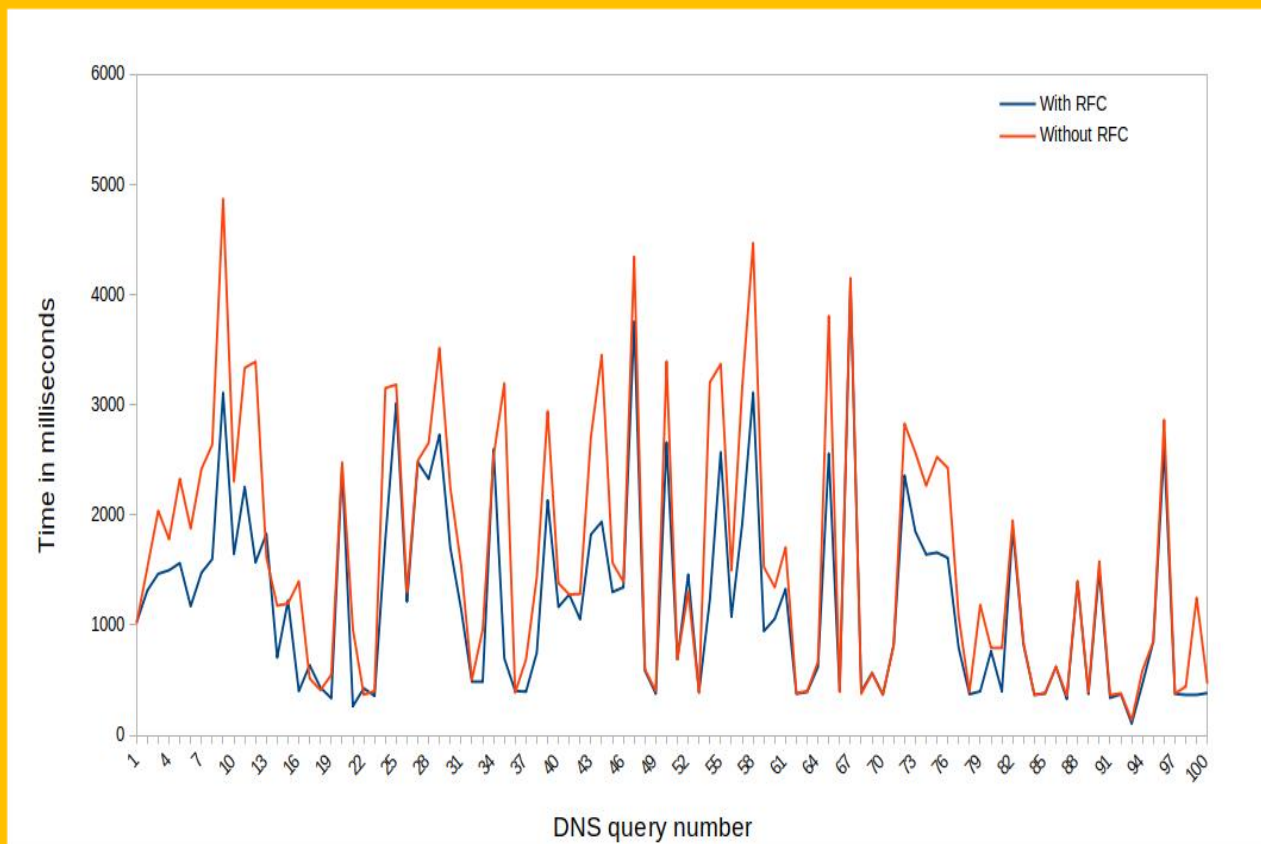


Figure.4: Comparison of domain resolution RTT of with and without RFC 7706

Conclusion

Using our set-up of open-public RR and querying it from terminal with public IP address helped us get genuine RTT for domain name resolution. By comparing the RTT obtained by querying the RR (*with and without RFC 7706*) with a dataset of 100 domain names, we can clearly understand that implementing RFC 7706, to run an instance of the root zone data in the loopback of RR, saves time, improves throughput and beneficial to the Internet ecosystem.

References:

- (1) RFC 7706: <https://tools.ietf.org/html/rfc7706>
- (2) Procedure for configuring DNS Bind server on CentOS 7:
<https://www.itzgeek.com/how-tos/linux/centos-how-tos/configure-dns-bind-server-on-centos-7-rhel-7.html>

Acknowledgements:

We express our sincere thanks to Internet Governance Division of [Ministry of Electronics & Information Technology \(MeitY\)](#) and [National Internet Exchange of India \(NIXI\)](#).

Cite as: Dr. Balaji Rajendran, Gopinath Palaniappan, Sanjay Adiwai, Shubham Goyal, Bindhumadhava B S (2020, January). Reducing RTT of DNS Query Resolution using RFC 7706. DOI: 10.17605/osf.io/t9n5x